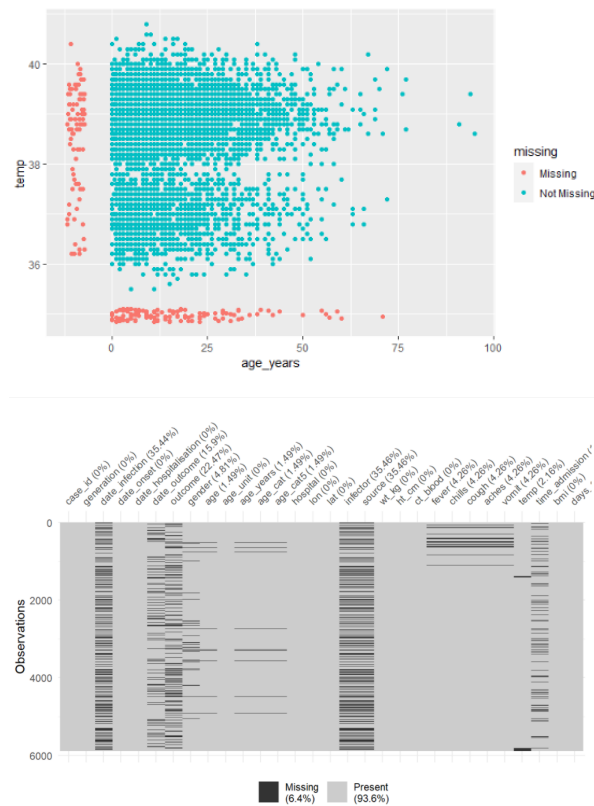


Missing data



This chapter will cover how to assess missingness, filter out rows by missingness, and handle how NA is displayed in plots.

Preparation

Load packages

This code chunk shows the loading of packages required for the analyses. In this handbook we emphasize `p_load()` from **pacman**, which installs the package if necessary *and* loads it for use. You can also load installed packages with `library()` from **base** R. See the page on [R basics] for more information on R packages.

```
pacman::p_load(
  rio,           # import/export
  tidyverse,    # data mgmt and viz
  naniar,       # assess and visualize missingness
```

```
mice          # missing data imputation
)
```

Installing package into 'C:/Users/Chris/AppData/Local/R/win-library/4.2'
(as 'lib' is unspecified)

Warning: unable to access index for repository <http://www.stats.ox.ac.uk/pub/RWin/bin/windows/contrib/4.2/PACKAGES>
cannot open URL 'http://www.stats.ox.ac.uk/pub/RWin/bin/windows/contrib/4.2/PACKAGES'

package 'mice' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
C:\Users\Chris\AppData\Local\Temp\RtmpgLoMfE\downloaded_packages

mice installed

Warning: package 'mice' was built under R version 4.2.3

Warning in pacman::p_load(rio, tidyverse, nanian, mice): Failed to install/load:
mice

Import data

We import the dataset of cases from a simulated Ebola epidemic. If you want to follow along, click to download the “clean” linelist (as .rds file). Import your data with the `import()` function from the **rio** package (it accepts many file types like .xlsx, .rds, .csv - see the [Import and export] page for details).

```
# import the linelist
linelist <- import("linelist_cleaned.rds")
```

The first 50 rows of the linelist are displayed below.

Show 5 entries

Search:

case_id	generation	date_infection	date_onset	date_hospitalisation	date_death
All	All	All	All	All	All
5fe599	4	2014-05-08	2014-05-13	2014-05-15	
8689b7	4		2014-05-13	2014-05-14	2014-05-15
11f8ea	2		2014-05-16	2014-05-18	2014-05-19
b8812a	3	2014-05-04	2014-05-18	2014-05-20	
893f25	3	2014-05-18	2014-05-21	2014-05-22	2014-05-23

Showing 1 to 5 of 50 entries

Previous

1

2

3

4

5

...

10

Next

Convert missing on import

When importing your data, be aware of values that should be classified as missing. For example, 99, 999, “Missing”, blank cells (“”), or cells with an empty space (“ ”). You can convert these to NA (R’s version of missing data) during the data import command.

See the page on importing page section on [Missing data](#) for details, as the exact syntax varies by file type.

Missing values in R

Below we explore ways that missingness is presented and assessed in R, along with some adjacent values and functions.

NA

In R, missing values are represented by a reserved (special) value - NA. Note that this is typed *without* quotes. “NA” is different and is just a normal character value (also a Beatles lyric from the song Hey Jude).

Your data may have other ways of representing missingness, such as “99”, or “Missing”, or “Unknown” - you may even have empty character value “ ” which looks “blank”, or a single space “ ”. Be aware of these and consider whether to [convert them to NA during import](#) or during data cleaning with `na_if()`.

In your data cleaning, you may also want to convert the other way - changing all NA to “Missing” or similar with `replace_na()` or with `fct_explicit_na()` for factors.

Versions of NA

Most of the time, NA represents a missing value and everything works fine. However, in some circumstances you may encounter the need for *variations* of NA specific to an object class (character, numeric, etc). This will be rare, but you should be aware.

The typical scenario for this is when creating a new column with the **dplyr** function `case_when()`. As described in the [Cleaning data and core functions](#) page, this function evaluates every row in the data frame, assess whether the rows meets specified logical criteria (right side of the code), and assigns the correct new value (left side of the code). *Importantly: all values on the right side must be the same class.*

```
linelist <- linelist %>%  
  
  # Create new "age_years" column from "age" column
```

```
mutate(age_years = case_when(
  age_unit == "years" ~ age,      # if age is given in years, assign original value
  age_unit == "months" ~ age/12,  # if age is given in months, divide by 12
  is.na(age_unit) ~ age,          # if age UNIT is missing, assume years
  TRUE ~ NA_real_)) # any other circumstance, assign missing
```

If you want NA on the right side, you may need to specify one of the special NA options listed below. If the other right side values are character, consider using “Missing” instead or otherwise use `NA_character_`. If they are all numeric, use `NA_real_`. If they are all dates or logical, you can use NA.

- NA - use for dates or logical TRUE/FALSE
- `NA_character_` - use for characters
- `NA_real_` - use for numeric

Again, it is not likely you will encounter these variations *unless* you are using `case_when()` to create a new column. See the [R documentation on NA](#) for more information.

NULL

NULL is another reserved value in R. It is the logical representation of a statement that is neither true nor false. It is returned by expressions or functions whose values are undefined. Generally do not assign NULL as a value, unless writing functions or perhaps writing a [\[shiny app\]](#)[Dashboards with Shiny] to return NULL in specific scenarios.

Null-ness can be assessed using `is.null()` and conversion can be made with `as.null()`.

See this [blog post](#) on the difference between NULL and NA.

NaN

Impossible values are represented by the special value NaN. An example of this is when you force R to divide 0 by 0. You can assess this with `is.nan()`. You may also encounter complementary functions including `is.infinite()` and `is.finite()`.

Inf

Inf represents an infinite value, such as when you divide a number by 0.

As an example of how this might impact your work: let’s say you have a vector/column `z` that contains these values: `z <- c(1, 22, NA, Inf, NaN, 5)`

If you want to use `max()` on the column to find the highest value, you can use the `na.rm = TRUE` to remove the NA from the calculation, but the `Inf` and `NaN` remain and `Inf` will be returned. To resolve this, you can use brackets `[]` and `is.finite()` to subset such that only finite values are used for the calculation: `max(z[is.finite(z)])`.

```
z <- c(1, 22, NA, Inf, NaN, 5)
max(z)                # returns NA
max(z, na.rm=T)       # returns Inf
max(z[is.finite(z)])  # returns 22
```

Examples

R command	Outcome
5 / 0	Inf
0 / 0	NaN
5 / NA	NA
5 / Inf 0NA - 5 NAInf / 5 Infclass(NA) "logical"class(NaN) "numeric"class(Inf) "numeric"class(NULL)‘	“NULL”

“NAs introduced by coercion” is a common warning message. This can happen if you attempt to make an illegal conversion like inserting a character value into a vector that is otherwise numeric.

```
as.numeric(c("10", "20", "thirty", "40"))
```

Warning: NAs introduced by coercion

```
[1] 10 20 NA 40
```

NULL is ignored in a vector.

```
my_vector <- c(25, NA, 10, NULL) # define
my_vector                # print
```

```
[1] 25 NA 10
```

Variance of one number results in NA.

```
var(22)
```

```
[1] NA
```

Useful functions

The following are useful **base** R functions when assessing or handling missing values:

`is.na()` and `!is.na()`

Use `is.na()` to identify missing values, or use its opposite (with `!` in front) to identify non-missing values. These both return a logical value (**TRUE** or **FALSE**). Remember that you can `sum()` the resulting vector to count the number **TRUE**, e.g. `sum(is.na(linelist$date_outcome))`.

```
my_vector <- c(1, 4, 56, NA, 5, NA, 22)
is.na(my_vector)
```

```
[1] FALSE FALSE FALSE  TRUE FALSE  TRUE FALSE
```

```
!is.na(my_vector)
```

```
[1]  TRUE  TRUE  TRUE FALSE  TRUE FALSE  TRUE
```

```
sum(is.na(my_vector))
```

```
[1] 2
```

`na.omit()`

This function, if applied to a data frame, will remove rows with *any* missing values. It is also from **base** R.

If applied to a vector, it will remove **NA** values from the vector it is applied to. For example:

```
na.omit(my_vector)
```

```
[1] 1 4 56 5 22
attr(,"na.action")
[1] 4 6
attr(,"class")
[1] "omit"
```

`drop_na()`

This is a **tidyr** function that is useful in a [data cleaning pipeline][Cleaning data and core functions]. If run with the parentheses empty, it removes rows with *any* missing values. If column names are specified in the parentheses, rows with missing values in those columns will be dropped. You can also use “tidyselect” syntax to specify the columns.

```
linelist %>%
  drop_na(case_id, date_onset, age) # drops rows missing values for any of these columns
```

`na.rm = TRUE`

When you run a mathematical function such as `max()`, `min()`, `sum()` or `mean()`, if there are any NA values present the returned value will be NA. This default behavior is intentional, so that you are alerted if any of your data are missing.

You can avoid this by removing missing values from the calculation. To do this, include the argument `na.rm = TRUE` (“na.rm” stands for “remove NA”).

```
my_vector <- c(1, 4, 56, NA, 5, NA, 22)

mean(my_vector)
```

```
[1] NA
```

```
mean(my_vector, na.rm = TRUE)
```

```
[1] 17.6
```


Using data with missing values

Filter out rows with missing values

To quickly remove rows with missing values, use the **dplyr** function `drop_na()`.

The original `linelist` has `nrow(linelist)` rows. The adjusted number of rows is shown below:

```
linelist %>%  
  drop_na() %>%      # remove rows with ANY missing values  
  nrow()
```

```
[1] 1818
```

You can specify to drop rows with missingness in certain columns:

```
linelist %>%  
  drop_na(date_onset) %>% # remove rows missing date_onset  
  nrow()
```

```
[1] 5632
```

You can list columns one after the other, or use “tidyselect” helper functions:

```
linelist %>%  
  drop_na(contains("date")) %>% # remove rows missing values in any "date" column  
  nrow()
```

```
[1] 3029
```

Handling NA in `ggplot()`

It is often wise to report the number of values excluded from a plot in a caption. Below is an example:

In `ggplot()`, you can add `labs()` and within it a `caption =`. In the caption, you can use `str_glue()` from **stringr** package to paste values together into a sentence dynamically so they will adjust to the data. An example is below:

- Note the use of `\n` for a new line.
- Note that if multiple column would contribute to values not being plotted (e.g. age or sex if those are reflected in the plot), then you must filter on those columns as well to correctly calculate the number not shown.

```
labs(
  title = "",
  y = "",
  x = "",
  caption = stringr::str_glue(
    "n = {nrow(central_data)} from Central Hospital;
    {nrow(central_data %>% filter(is.na(date_onset)))} cases missing date of onset and not s
```

Sometimes, it can be easier to save the string as an object in commands prior to the `ggplot()` command, and simply reference the named string object within the `str_glue()`.