Recoding Values

Here are a few scenarios where you need to re-code (change) values:

- to edit one specific value (e.g. one date with an incorrect year or format)

- to reconcile values not spelled the same
- to create a new column of categorical values

- to create a new column of numeric categories (e.g. age categories)

> **Data Preparation**

Manually changing specific values

To change values manually you can use the `recode()` function within the `mutate()` function.

Imagine there is a nonsensical date in the data (e.g. "2014-14-15"): you could fix the date manually in the raw source data, or, you could write the change into the cleaning pipeline via `mutate()` and `recode()`. The latter is more transparent and reproducible to anyone else seeking to understand or repeat your analysis.

```
# fix incorrect values                # old value      # new value
linelist <- linelist %>%
  mutate(date_onset = recode(date_onset, "2014-14-15" = "2014-04-15"))
```

The `mutate()` line above can be read as: "mutate the column `date_onset` to equal the column `date_onset` re-coded so that OLD VALUE is changed to NEW VALUE". Note that this pattern (OLD = NEW) for `recode()` is the opposite of most R patterns (new = old). The R development community is working on revising this.

By logic

Below we demonstrate how to re-code values in a column using logic and conditions:

- Using `replace()`, `ifelse()` and `if_else()` for simple logic
- Using `case_when()` for more complex logic

Simple logic

* replace()

To re-code with simple logical criteria, you can use `replace()` within `mutate()`. `replace()` is a function from **base** R. Use a logic condition to specify the rows to change . The general syntax is:

`mutate(col_to_change = replace(col_to_change, criteria for rows, new value))`.

One common situation to use `replace()` is **changing just one value in one row, using an unique row identifier**. Below, the gender is changed to "Female" in the row where the column `case_id` is "2195".

```
# Example: change gender of one specific observation to "Female"
linelist <- linelist %>%
  mutate(gender = replace(gender, case_id == "2195", "Female"))
```

The equivalent command using **base** R syntax and indexing brackets `[ ]` is below. It reads as "Change the value of the dataframe `linelist`'s column `gender` (for the rows where `linelist`'s column `case_id` has the value '2195') to 'Female' ".

```
linelist$gender[linelist$case_id == "2195"] <- "Female"
```


* ifelse() and if_else()

Another tool for simple logic is `ifelse()` and its partner `if_else()`. However, in most cases for re-coding it is more clear to use `case_when()` (detailed below). These "if else" commands are simplified versions of an `if` and `else` programming statement. The general syntax is:
`ifelse(condition, value to return if condition evaluates to TRUE, value to return if condition evaluates to FALSE)`
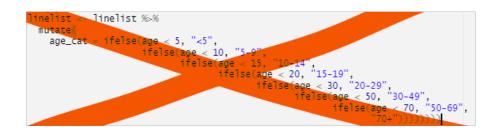
Below, the column `source_known` is defined. Its value in a given row is set to "known" if the row's value in column `source` is *not* missing. If the value in `source` *is* missing, then the value in `source_known` is set to "unknown".

```
linelist <- linelist %>%
  mutate(source_known = ifelse(!is.na(source), "known", "unknown"))
```

`if_else()` is a special version from **dplyr** that handles dates. Note that if the 'true' value is a date, the 'false' value must also qualify a date, hence using the special value `NA_real_` instead of just `NA`.

2

```
# Create a date of death column, which is NA if patient has not died.
linelist <- linelist %>%
  mutate(date_death = if_else(outcome == "Death", date_outcome, NA_real_))
```

Avoid stringing together many ifelse commands... use `case_when()` instead!
`case_when()` is much easier to read and you'll make fewer errors.

Outside of the context of a data frame, if you want to have an object used in your code switch its value, consider using `switch()` from **base** R.

Complex logic

Use **dplyr**'s `case_when()` if you are re-coding into many new groups, or if you need to use complex logic statements to re-code values. This function evaluates every row in the data frame, assess whether the rows meets specified criteria, and assigns the correct new value.

`case_when()` commands consist of statements that have a Right-Hand Side (RHS) and a Left-Hand Side (LHS) separated by a "tilde" `~`. The logic criteria are in the left side and the pursuant values are in the right side of each statement. Statements are separated by commas.

For example, here we utilize the columns `age` and `age_unit` to create a column `age_years`:

```
linelist <- linelist %>%
  mutate(age_years = case_when(
      age_unit == "years"  ~ age,        # if age unit is years
      age_unit == "months" ~ age/12,     # if age unit is months, divide age by 12
      is.na(age_unit)      ~ age))       # if age unit is missing, assume years
                                         # any other circumstance, assign NA (missing)
```

As each row in the data is evaluated, the criteria are applied/evaluated in the order the `case_when()` statements are written - from top-to-bottom. If the top criteria evaluates to `TRUE` for a given row, the RHS value is assigned, and the remaining criteria are not even tested for that row in the data. Thus, it is best to write the most specific criteria first, and the most general last. A data row that does not meet any of the RHS criteria will be assigned `NA`.

Sometimes, you may with to write a final statement that assigns a value for all other scenarios not described by one of the previous lines. To do this, place `TRUE` on the left-side, which will capture any row that did not meet any of the previous criteria. The right-side of this statement could be assigned a value like "check me!" or missing.

Below is another example of `case_when()` used to create a new column with the patient classification, according to a case definition for confirmed and suspect cases:

```r
linelist <- linelist %>%
     mutate(case_status = case_when(

          # if patient had lab test and it is positive,
          # then they are marked as a confirmed case
          ct_blood < 20                     ~ "Confirmed",

          # given that a patient does not have a positive lab result,
          # if patient has a "source" (epidemiological link) AND has fever,
          # then they are marked as a suspect case
          !is.na(source) & fever == "yes" ~ "Suspect",

          # any other patient not addressed above
          # is marked for follow up
          TRUE                              ~ "To investigate"))
```

Numeric categories

Here we describe some special approaches for creating categories from numerical columns. Common examples include age categories, groups of lab values, etc. Here we will discuss:

- `age_categories()`, from the **epikit** package

- `cut()`, from **base** R

- `case_when()`

- quantile breaks with `quantile()` and `ntile()`

Review distribution

For this example we will create an `age_cat` column using the `age_years` column.

```
#check the class of the linelist variable age
class(linelist$age_years)
```

```
[1] "numeric"
```

`cut()`

The basic syntax within `cut()` is to first provide the numeric column to be cut (`age_years`), and then the *breaks* argument, which is a numeric vector `c()` of break points. Using `cut()`, the resulting column is an ordered factor.

By default, the categorization occurs so that the right/upper side is "open" and inclusive (and the left/lower side is "closed" or exclusive). This is the opposite behavior from the `age_categories()` function. The default labels use the notation "(A, B]", which means A is not included but B is. **Reverse this behavior by providing the `right = TRUE` argument**.

Thus, by default, "0" values are excluded from the lowest group, and categorized as `NA`! "0" values could be infants coded as age 0 so be careful! To change this, add the argument `include.lowest = TRUE` so that any "0" values will be included in the lowest group. The automatically-generated label for the lowest category will then be "[A],B]". Note that if you include the `include.lowest = TRUE` argument **and** `right = TRUE`, the extreme inclusion will now apply to the *highest* break point value and category, not the lowest.

You can provide a vector of customized labels using the `labels =` argument. As these are manually written, be very careful to ensure they are accurate! Check your work using cross-tabulation, as described below.

An example of `cut()` applied to `age_years` to make the new variable `age_cat` is below:

```
# Create new variable, by cutting the numeric age variable
# lower break is excluded but upper break is included in each category
linelist <- linelist %>%
  mutate(
    age_cat = cut(
      age_years,
      breaks = c(0, 5, 10, 15, 20,
                 30, 50, 70, 100),
      include.lowest = TRUE       # include 0 in lowest group
```

5

```
    ))

# tabulate the number of observations per group
table(linelist$age_cat, useNA = "always")
```

```
  [0,5]   (5,10]  (10,15]  (15,20]  (20,30]  (30,50]  (50,70] (70,100]
   1315     1065      930      696     1013      694       84        5
  <NA>
    86
```

Check your work!!! Verify that each age value was assigned to the correct category by cross-tabulating the numeric and category columns. Examine assignment of boundary values (e.g. 15, if neighboring categories are 10-15 and 16-20).

Quantile breaks

In common understanding, "quantiles" or "percentiles" typically refer to a value below which a proportion of values fall. For example, the 95th percentile of ages in `linelist` would be the age below which 95% of the age fall.

However in common speech, "quartiles" and "deciles" can also refer to the *groups of data* as equally divided into 4, or 10 groups (note there will be one more break point than group).

To get quantile break points, you can use `quantile()` from the **stats** package from **base** R. You provide a numeric vector (e.g. a column in a dataset) and vector of numeric probability values ranging from 0 to 1.0. The break points are returned as a numeric vector. Explore the details of the statistical methodologies by entering `?quantile`.

- If your input numeric vector has any missing values it is best to set `na.rm = TRUE`

- Set `names = FALSE` to get an un-named numeric vector

```
quantile(linelist$age_years,                  # specify numeric vector to work on
  probs = c(0, .25, .50, .75, .90, .95),      # specify the percentiles you want
  na.rm = TRUE)                               # ignore missing values
```

```
  0%   25%   50%   75%   90%   95%
 0.0   6.0  13.0  23.0  33.9  41.0
```

You can use the results of `quantile()` as break points in `age_categories()` or `cut()`. Below we create a new column `deciles` using `cut()` where the breaks are defined using `quantiles()` on `age_years`. Below, we display the results using `tabyl()` from **janitor** so you can see the percentages. Note how they are not exactly 10% in each group.

```
linelist %>%                            # begin with linelist
  mutate(deciles = cut(age_years,       # create new column decile as cut() on column
    breaks = quantile(                  # define cut breaks using quantile()
      age_years,                        # operate on age_years
      probs = seq(0, 1, by = 0.1),      # 0.0 to 1.0 by 0.1
      na.rm = TRUE),                    # ignore missing values
    include.lowest = TRUE)) %>%         # for cut() include age 0
  janitor::tabyl(deciles)              # pipe to table to display
```

```
   deciles   n   percent valid_percent
     [0,2] 658 0.11175272    0.11340917
     (2,5] 657 0.11158288    0.11323681
     (5,7] 447 0.07591712    0.07704240
    (7,10] 618 0.10495924    0.10651499
   (10,13] 572 0.09714674    0.09858669
   (13,17] 674 0.11447011    0.11616684
   (17,21] 520 0.08831522    0.08962427
   (21,26] 547 0.09290082    0.09427784
 (26,33.9] 528 0.08967391    0.09100310
 (33.9,84] 581 0.09867527    0.10013788
      <NA>  86 0.01460598            NA
```